

DOC2VEC

[] stem_df.head(2)

	0_left	1	2	3	4	5	6	7	8	9	...	91	92	93	94	95	96	97	98	99	0_right
0	-0.007656	-0.017189	0.058634	0.072262	-0.014387	0.013050	-0.043178	-0.065979	0.055877	-0.036416	...	-0.054352	0.042411	-0.017314	-0.025635	0.013351	-0.065088	0.013830	0.062138	0.054103	0
1	-0.031230	0.023416	0.008874	0.047421	-0.016033	-0.025616	-0.050069	-0.018542	0.036259	-0.002108	...	-0.076787	0.021046	-0.026546	0.046404	-0.064159	0.009849	0.009103	0.022082	0.063062	1

2 rows × 201 columns

[30] df_lemma.head(2)

	q1	q2	is_duplicate
48958	job microsoft certification	microsoft certification like mose mc aid get L...	0
77587	write article art gallery	write art easy	0

[] lemma_df.head(2)

	0_left	1	2	3	4	5	6	7	8	9	...	41	42	43	44	45	46	47	48	49	0_right
0	-0.250281	0.268303	-0.507923	-0.090775	-0.015231	-0.105132	-0.230446	0.054211	0.073804	-0.101047	...	0.015447	-0.113064	0.021469	0.100241	-0.08906	0.092462	0.195167	0.278540	-0.025395	0
1	0.042417	0.109407	0.018437	-0.006771	-0.064517	-0.041608	-0.039554	-0.070556	0.149302	-0.058216	...	0.078449	0.072881	-0.073291	-0.031450	0.07713	0.034200	0.182466	0.055447	0.062686	0

2 rows × 101 columns

RANDOM FOREST

[5] #Random Forest

```
rf_clf = RandomForestClassifier(random_state=0)
rf_param_grid = {'max_depth': [20, None],
                 'min_samples_split': [20]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resource=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, 'f1acc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

[75] #Random Forest

```
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)
```

RandomForestClassifier(max_depth=20, min_samples_split=20, n_estimators=5, random_state=0)

Classification Report - RF

	precision	recall	f1-score	support
0	1.00	1.00	1.00	4748
1	0.99	1.00	1.00	2752
accuracy	1.00	1.00	1.00	7500
macro avg	1.00	1.00	1.00	7500
weighted avg	1.00	1.00	1.00	7500

Confusion Matrix - RF

```
[[4727 21]
 [ 1 2751]]
```

Executing [54]: Cell @ f0 : _run_search() : evaluate_candidates() : _call_() : retrieve() : wrap_future_result() : result() : wait()

[5] #Random Forest

```
rf_clf = RandomForestClassifier(random_state=0)
rf_param_grid = {'max_depth': [5],
                 'min_samples_split': [20]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resource=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, 'f1acc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

[75] #Random Forest

```
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)
```

RandomForestClassifier(max_depth=5, min_samples_split=20, n_estimators=5, random_state=0)

Classification Report - RF

	precision	recall	f1-score	support
0	1.00	0.59	0.82	6883
1	0.25	1.00	0.40	697
accuracy	0.63	0.85	0.72	7580
macro avg	0.63	0.85	0.61	7580
weighted avg	0.93	0.72	0.78	7580

Confusion Matrix - RF

```
[[4728 2075]
 [ 0 697]]
```

6s - completed at 14:54

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = { "max_depth": [5],
                  "min_samples_split": [100]
                }

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                               max_resources=5,
                               random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, 'f1_score', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

```
#Random Forest
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)
```

RandomForestClassifier(max_depth=5, min_samples_split=100, n_estimators=5, random_state=0)

Classification Report - RF

	precision	recall	f1-score	support
0	1.00	0.69	0.81	6890
1	0.22	1.00	0.36	610
accuracy	0.51	0.54	0.71	7500
macro avg	0.54	0.59	0.59	7500
weighted avg	0.54	0.71	0.78	7500

Confusion Matrix - RF

```
[[4728 2162]
 [  0  610]]
```

10s completed at 14:55

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = { "max_depth": [4],
                  "min_samples_split": [100]
                }

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                               max_resources=5,
                               random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, 'f1_score', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

```
#Random Forest
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)
```

RandomForestClassifier(max_depth=4, min_samples_split=100, n_estimators=5, random_state=0)

Classification Report - RF

	precision	recall	f1-score	support
0	1.00	0.88	0.94	5381
1	0.76	1.00	0.87	2119
accuracy	0.88	0.54	0.91	7500
macro avg	0.88	0.54	0.90	7500
weighted avg	0.93	0.91	0.92	7500

Confusion Matrix - RF

```
[[4728  653]
 [  0 2119]]
```

4s completed at 14:57

SELECTED HP

```
def rf_clf(X_train, X_test, y_train, y_test):
    """Parameter tuning for Random Forest Classifier and model fit"""
    #split = ShuffleSplit(n_splits = 1, test_size = .1, random_state = 0)
    rf_clf = RandomForestClassifier(random_state = 0)
    rf_param_grid = { "max_depth": [1, 100],
                     "min_samples_split": [40]
                   }

    rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators', max_resources=5,
                                   random_state=0).fit(X_train, y_train)

    rf_model = rf_search.best_estimator_
    y_pred = rf_model.predict(X_test)

    log_loss = metrics.log_loss(y_test, y_pred)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    precision = metrics.precision_score(y_test, y_pred)
    recall = metrics.recall_score(y_test, y_pred)
    f1_score = metrics.f1_score(y_test, y_pred)
    print(rf_model)
    #print('log_loss', rf_loss_bow, 'f1_score', rf_acc_bow)
    return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

SVC

```
splits = ShuffleSplit(n_splits = 1, test_size = .2, random_state = 0)

#svc_param_grid = {'C': [0.1, 10, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'linear', 'poly', 'rbf']}
#svc_param_grid = {'C': [0.1, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'poly', 'rbf']}
svc_param_grid = {'C': [10],
                  'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                  'kernel': ['sigmoid']}
svc_clf = SVC()
# gamma='scale'
svc_search = HalvingGridSearchCV(svc_clf, svc_param_grid, factor = 2, scoring = 'accuracy')
svc_search.fit(X_train, y_train)

svc_model = svc_search.best_params_
svc_model_est = svc_search.best_estimator_

y_pred = svc_model_est.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred, average='weighted')
recall = metrics.recall_score(y_test, y_pred, average='weighted')
f1_score = metrics.f1_score(y_test, y_pred, average='macro')
print(svc_model)
return accuracy, precision, recall, f1_score, log_loss, y_pred, svc_model_est
```

```
# SVC - using param grid - 100, 30, 1
acc_svc_100ma, precision_svc_100ma, recall_svc_100ma, f1_score_svc_100ma, log_loss_svc_100ma, y_pred_svc_100ma, svc_model_100ma = svc_clf(X_train_100ma, X_test_100ma, y_train_100ma, y_test_100ma)

report_svc_100ma = classification_report(y_pred_svc_100ma, y_test_100ma)
matrix_svc_100ma = confusion_matrix(y_pred_svc_100ma, y_test_100ma)

print('Classification Report - SVC\n', report_svc_100ma)
print()
print('Confusion Matrix - SVC\n', matrix_svc_100ma)
```

```
{'C': 1, 'gamma': 0.01, 'kernel': 'rbf'}
Classification Report - SVC
      precision    recall  f1-score   support

0         1.00      1.00      1.00        4733
1         1.00      1.00      1.00        2767

accuracy          1.00      1.00      1.00       7500
macro avg          1.00      1.00      1.00       7500
weighted avg       1.00      1.00      1.00       7500
```

```
Confusion Matrix - SVC
[[4733  0]
 [ 0 2767]]
```

1m 10s completed at 14:50

SELECTED

```
def svc_clf(X_train, X_test, y_train, y_test):
    """Parameter tuning for svm classifier and model fit"""

    splits = ShuffleSplit(n_splits = 1, test_size = .2, random_state = 0)

    #svc_param_grid = {'C': [0.1, 10, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'linear', 'poly', 'rbf']}
    #svc_param_grid = {'C': [0.1, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'poly', 'rbf']}
    svc_param_grid = {'C': [10],
                      'gamma': [1],
                      'kernel': ['sigmoid']}
    svc_clf = SVC()
    # gamma='scale'
    svc_search = HalvingGridSearchCV(svc_clf, svc_param_grid, factor = 2, scoring = 'accuracy')
    svc_search.fit(X_train, y_train)

    svc_model = svc_search.best_params_
    svc_model_est = svc_search.best_estimator_

    y_pred = svc_model_est.predict(X_test)

    log_loss = metrics.log_loss(y_test, y_pred)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    precision = metrics.precision_score(y_test, y_pred, average='weighted')
    recall = metrics.recall_score(y_test, y_pred, average='weighted')
    f1_score = metrics.f1_score(y_test, y_pred, average='macro')
    print(svc_model)
    return accuracy, precision, recall, f1_score, log_loss, y_pred, svc_model_est
```

DOC2VEC – RANDOM FOREST

VECTOR 20

```
#Random forest # doc0, window=1, vector=20, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       1.00      1.00      1.00        4737
    1       1.00      1.00      1.00        2763

 accuracy          1.00          1.00          1.00       7500
 macro avg          1.00          1.00          1.00       7500
 weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - RF
[[4732  5]
 [ 1 2762]]

[149] #Random forest # doc1, window=1, vector=20, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=20, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       1.00      1.00      1.00        4746
    1       0.99      1.00      1.00        2754

 accuracy          1.00          1.00          1.00       7500
 macro avg          1.00          1.00          1.00       7500
 weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - RF
[[4732 14]
 [ 1 2753]]
```

```
[ ] #Random forest # doc0, window=1, vector=100, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=20, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.99      0.98      0.98        4785
    1       0.96      0.98      0.97        2715

 accuracy          0.97          0.98          0.98       7500
 macro avg          0.97          0.98          0.98       7500
 weighted avg          0.98          0.98          0.98       7500

Confusion Matrix - RF
[[4674 111]
 [ 59 2656]]

[ ] #Random forest # doc0, window=1, vector=20, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       1.00      1.00      1.00        4737
    1       1.00      1.00      1.00        2763

 accuracy          1.00          1.00          1.00       7500
 macro avg          1.00          1.00          1.00       7500
 weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - RF
[[4732  5]
 [ 1 2762]]
```

DM = 1 - VECTOR=50 – WINDOW= 3 - EPOCHS=30

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {"max_depth": [10, None],
                 "min_samples_split": [2, 4]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy',
                                max_resources=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print(log_loss, rf_loss_bow, "\nacc", rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=4, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.99      0.98      0.98      4763
1           0.96      0.98      0.97      2737

accuracy          0.98      0.98      0.98      7500
macro avg         0.97      0.98      0.98      7500
weighted avg      0.98      0.98      0.98      7500

Confusion Matrix - RF
[[4661 182]
 [ 62 2670]]
```

21s completed at 11:16

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {"max_depth": [10, None],
                 "min_samples_split": [20, 40]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy',
                                max_resources=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print(log_loss, rf_loss_bow, "\nacc", rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.98      0.95      0.97      4880
1           0.92      0.97      0.94      2620

accuracy          0.96      0.96      0.96      7500
macro avg         0.95      0.96      0.95      7500
weighted avg      0.96      0.96      0.96      7500

Confusion Matrix - RF
[[4646 234]
 [ 82 2538]]
```

16s completed at 11:17

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {'max_depth': [40, None],
                 "min_samples_split": [20, 40]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resource=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_

y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, '\nacc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.98      0.95      0.97      4880
    1       0.92      0.97      0.94      2620

 accuracy      0.96      0.96      0.96      7500
 macro avg     0.95      0.96      0.95      7500
 weighted avg  0.96      0.96      0.96      7500

Confusion Matrix - RF
[[4646  234]
 [  82 2538]]
```

15s completed at 11:18

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {'max_depth': [40, None],
                 "min_samples_split": [20, 40]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resource=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_

y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, '\nacc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(max_depth=40, min_samples_split=20, n_estimators=3,
                        random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.99      0.98      0.99      4789
    1       0.97      0.99      0.98      2711

 accuracy      0.98      0.98      0.98      7500
 macro avg     0.98      0.98      0.98      7500
 weighted avg  0.98      0.98      0.98      7500

Confusion Matrix - RF
[[4696   93]
 [  32 2679]]
```

16s completed at 11:19

VECTOR 100 VS 200 (DM = 1 – WINDOW = 3, EPOCHS = 30)

```
#Random Forest # dm1, window3, vector=200, epochs=30
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, bcrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.96     0.91     0.94     5002
1           0.84     0.93     0.89     2498

accuracy          0.90     0.92     0.91     7500
macro avg         0.90     0.92     0.91     7500
weighted avg      0.92     0.92     0.92     7500

Confusion Matrix - RF
[[4565 437]
 [108 2330]]
```

CHANGING NUMBER OF WINDOW

```
[37] #Random Forest # 100 - 30 - 1
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, bcrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.99     0.96     0.98     4865
1           0.93     0.98     0.96     2635

accuracy          0.96     0.97     0.97     7500
macro avg         0.96     0.97     0.97     7500
weighted avg      0.97     0.97     0.97     7500

Confusion Matrix - RF
[[4680 185]
 [ 53 2520]]
```

```
#Random Forest # 100 - 30 - 3
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, bcrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=20, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.96     0.93     0.95     4867
1           0.88     0.93     0.91     2633

accuracy          0.92     0.93     0.93     7500
macro avg         0.92     0.93     0.93     7500
weighted avg      0.93     0.93     0.93     7500

Confusion Matrix - RF
[[4548 319]
 [191 2442]]
```

```
#Random Forest # 100 - 30 - 5
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, bcrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.99     0.95     0.97     4915
1           0.91     0.98     0.95     2585

accuracy          0.95     0.97     0.96     7500
macro avg         0.95     0.97     0.96     7500
weighted avg      0.96     0.96     0.96     7500

Confusion Matrix - RF
[[4679 236]
 [ 54 2521]]
```

RESULTS

Lemmatized without stopwords

Not tokenized – 100 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

0       1.00      0.50      0.50       150
1       0.93      1.00      0.96       104

accuracy    0.95
macro avg   0.95      0.95      0.97       300
weighted avg 0.95      0.97      0.97       300

Confusion Matrix - RF
[[150  0]
 [ 0 104]]
{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

0       0.76      0.77      0.77       185
1       0.62      0.61      0.62       115

accuracy    0.69
macro avg   0.69      0.69      0.71       300
weighted avg 0.71      0.71      0.71       300

Confusion Matrix - SVC
[[145  40]
 [ 40 155]]
```

Tokenized – 100 VECTORS

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

0       0.97      0.94      0.96       485
1       0.90      0.95      0.92       265

accuracy    0.94
macro avg   0.94      0.95      0.94       750
weighted avg 0.95      0.95      0.95       750

Confusion Matrix - RF
[[485  0]
 [ 0 265]]
{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

0       0.78      0.76      0.77       485
1       0.58      0.61      0.59       265

accuracy    0.68
macro avg   0.68      0.68      0.71       750
weighted avg 0.71      0.71      0.71       750

Confusion Matrix - SVC
[[367 118]
 [108 192]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

0       0.88      0.88      0.88       475
1       0.79      0.81      0.80       275

accuracy    0.84
macro avg   0.84      0.85      0.85       750
weighted avg 0.85      0.85      0.85       750

Confusion Matrix - RF
[[475  0]
 [ 0 275]]
{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

0       0.79      0.78      0.78       475
1       0.62      0.64      0.63       272

accuracy    0.78
macro avg   0.78      0.71      0.75       750
weighted avg 0.73      0.73      0.73       750

Confusion Matrix - SVC
[[371 104]
 [ 99 179]]
```


Lemmatized with stopwords

Not tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

 0      1.00    0.97    0.99    103
 1      0.96    1.00    0.98    107

 accuracy      0.98
 macro avg     0.98
 weighted avg   0.98

Confusion Matrix - RF
[[100  3]
 [ 0 107]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

 0      0.74    0.74    0.74    103
 1      0.96    0.97    0.97    111

 accuracy      0.85
 macro avg     0.85
 weighted avg   0.88

Confusion Matrix - SVC
[[100  3]
 [ 0 111]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

 0      0.97    0.92    0.95    495
 1      0.86    0.95    0.90    255

 accuracy      0.91
 macro avg     0.91
 weighted avg   0.93

Confusion Matrix - RF
[[486  9]
 [ 0 261]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

 0      0.72    0.71    0.72    472
 1      0.92    0.92    0.92    278

 accuracy      0.84
 macro avg     0.82
 weighted avg   0.84

Confusion Matrix - SVC
[[337 135]
 [133 141]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

 0      0.94    0.89    0.90    513
 1      0.79    0.88    0.83    237

 accuracy      0.87
 macro avg     0.84
 weighted avg   0.88

Confusion Matrix - RF
[[442  71]
 [ 0 206]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

 0      0.75    0.74    0.74    475
 1      0.96    0.97    0.96    275

 accuracy      0.88
 macro avg     0.86
 weighted avg   0.88

Confusion Matrix - SVC
[[383 194]
 [129 186]]
```

With stopwords, not lemmatized nor stemmed

Not tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.99      0.98      0.98       193
     1       0.94      0.96      0.95       257

 accuracy      0.96
 macro avg     0.96      0.97      0.97       300
 weighted avg   0.97      0.97      0.97       300

Confusion Matrix - RF
[[186  2]
 [ 3 267]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.77      0.77      0.77       187
     1       0.62      0.61      0.61       113

 accuracy      0.69
 macro avg     0.69      0.69      0.69       300
 weighted avg   0.71      0.71      0.71       300

Confusion Matrix - SVC
[[144 43]
 [ 44 69]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.97      0.96      0.96       483
     1       0.96      0.98      0.97       267

 accuracy      0.96
 macro avg     0.96      0.96      0.96       750
 weighted avg   0.96      0.96      0.96       750

Confusion Matrix - RF
[[465 18]
 [ 15 282]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.69      0.69      0.69       478
     1       0.49      0.49      0.49       272

 accuracy      0.62
 macro avg     0.69      0.59      0.59       750
 weighted avg   0.62      0.62      0.62       750

Confusion Matrix - SVC
[[106 144]
 [144 196]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.97      0.93      0.95       467
     1       0.88      0.94      0.91       283

 accuracy      0.93
 macro avg     0.93      0.94      0.93       750
 weighted avg   0.94      0.93      0.94       750

Confusion Matrix - RF
[[454 13]
 [ 31 297]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.70      0.69      0.70       473
     1       0.48      0.50      0.49       277

 accuracy      0.69
 macro avg     0.69      0.59      0.59       750
 weighted avg   0.62      0.62      0.62       750

Confusion Matrix - SVC
[[120 145]
 [142 193]]
```

Without stopwords, not lemmatized nor stemmed

No tokenized – 100 vectors

```
RandomForestClassifier(max_depth=6, min_samples_split=4, n_estimators=5, random_state=4)
Classification Report - RF
      precision    recall  f1-score   support

     0       1.00      0.84      0.91      223
     1       0.49      1.00      0.61       77

 accuracy      0.84      0.92      0.88      300
 macro avg      0.74      0.92      0.86      300
 weighted avg      0.92      0.88      0.89      300

Confusion Matrix - RF
[[188  35]
 [  9  77]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.82      0.75      0.80      188
     1       0.62      0.68      0.64      112

 accuracy      0.73      0.73      0.73      300
 macro avg      0.73      0.73      0.73      300
 weighted avg      0.75      0.75      0.75      300

Confusion Matrix - SVC
[[155  43]
 [ 33  89]]
```

Tokenized - vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=4)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.97      0.97      0.97      471
     1       0.95      0.96      0.96      279

 accuracy      0.97      0.97      0.97      750
 macro avg      0.96      0.96      0.96      750
 weighted avg      0.97      0.97      0.97      750

Confusion Matrix - RF
[[466  25]
 [ 25  269]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.79      0.78      0.78      474
     1       0.62      0.63      0.63      276

 accuracy      0.71      0.72      0.72      750
 macro avg      0.71      0.71      0.71      750
 weighted avg      0.73      0.73      0.73      750

Confusion Matrix - SVC
[[369 186]
 [186 171]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=4)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.96      0.98      0.97      468
     1       0.82      0.92      0.87      282

 accuracy      0.89      0.95      0.92      750
 macro avg      0.89      0.95      0.92      750
 weighted avg      0.91      0.95      0.93      750

Confusion Matrix - RF
[[461  47]
 [ 47  213]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.81      0.79      0.80      478
     1       0.65      0.67      0.66      272

 accuracy      0.73      0.73      0.73      750
 macro avg      0.73      0.73      0.73      750
 weighted avg      0.75      0.75      0.75      750

Confusion Matrix - SVC
[[379  99]
 [ 99 181]]
```

Stemmed with stopwords

Not tokenized – 100 vectors

```
RandomForestClassifier(m_n_samples_split=40, n_estimators=5, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0       0.96       0.91       0.93       180
1       0.84       0.70       0.76       182

accuracy    0.91       0.91       0.91       360
macro avg   0.90       0.82       0.85       360
weighted avg 0.92       0.81       0.85       360

Confusion Matrix - RF
[[180  0]
 [ 0 182]]

{'C': 18, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
precision    recall  f1-score   support

0       0.91       0.85       0.88       203
1       0.72       0.84       0.78       97

accuracy    0.86       0.84       0.84       300
macro avg   0.82       0.84       0.83       300
weighted avg 0.85       0.84       0.85       300

Confusion Matrix - SVC
[[172  31]
 [ 36  111]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(m_n_samples_split=40, n_estimators=5, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0       0.98       0.95       0.97       488
1       0.91       0.97       0.94       264

accuracy    0.95       0.96       0.96       752
macro avg   0.95       0.96       0.96       752
weighted avg 0.96       0.96       0.96       752

Confusion Matrix - RF
[[482  64]
 [  0 264]]

{'C': 18, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
precision    recall  f1-score   support

0       0.75       0.72       0.74       482
1       0.53       0.56       0.54       264

accuracy    0.67       0.64       0.67       752
macro avg   0.64       0.64       0.64       752
weighted avg 0.67       0.67       0.67       752

Confusion Matrix - SVC
[[391  91]
 [119 169]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(m_n_samples_split=40, n_estimators=5, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0       0.94       0.87       0.91       518
1       0.76       0.89       0.82       240

accuracy    0.86       0.88       0.86       758
macro avg   0.85       0.88       0.86       758
weighted avg 0.89       0.88       0.88       758

Confusion Matrix - RF
[[484  34]
 [ 36 244]]

{'C': 18, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
precision    recall  f1-score   support

0       0.75       0.74       0.75       482
1       0.56       0.59       0.57       268

accuracy    0.66       0.67       0.66       750
macro avg   0.66       0.67       0.66       750
weighted avg 0.69       0.69       0.69       750

Confusion Matrix - SVC
[[399  83]
 [151 167]]
```

Stemmed without stopwords

No Tokenized – 100 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.97      0.93      0.95        196
     1       0.88      0.95      0.92        184

 accuracy      0.93      0.94      0.94        380
 macro avg     0.93      0.94      0.93        380
 weighted avg  0.94      0.94      0.94        380

Confusion Matrix - RF
[[183  31]
 [ 9  91]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.95      0.85      0.90        200
     1       0.72      0.89      0.80        91

 accuracy      0.81      0.87      0.85        290
 macro avg     0.83      0.87      0.85        290
 weighted avg  0.85      0.86      0.87        290

Confusion Matrix - SVC
[[176  24]
 [ 18  65]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.98      0.97      0.98        471
     1       0.95      0.96      0.96        277

 accuracy      0.97      0.97      0.97        750
 macro avg     0.97      0.97      0.97        750
 weighted avg  0.97      0.97      0.97        750

Confusion Matrix - RF
[[466  25]
 [ 18  287]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.72      0.72      0.72        471
     1       0.53      0.53      0.53        279

 accuracy      0.62      0.63      0.63        750
 macro avg     0.62      0.63      0.63        750
 weighted avg  0.65      0.63      0.65        750

Confusion Matrix - SVC
[[339 132]
 [111 140]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.93      0.88      0.90        497
     1       0.78      0.86      0.82        253

 accuracy      0.85      0.87      0.86        750
 macro avg     0.85      0.87      0.86        750
 weighted avg  0.88      0.87      0.87        750

Confusion Matrix - RF
[[450  47]
 [ 16 248]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.71      0.72      0.71        462
     1       0.54      0.52      0.53        288

 accuracy      0.62      0.64      0.64        750
 macro avg     0.62      0.62      0.62        750
 weighted avg  0.64      0.64      0.64        750

Confusion Matrix - SVC
[[332 130]
 [126 140]]
```