

DOC2VEC

[] stem_df.head(2)

	0_left	1	2	3	4	5	6	7	8	9	...	91	92	93	94	95	96	97	98	99	0_right
0	-0.007656	-0.017189	0.058634	0.072262	-0.014387	0.013050	-0.043178	-0.065979	0.055877	-0.036416	...	-0.054352	0.042411	-0.017314	-0.025635	0.013351	-0.065088	0.013830	0.062138	0.054103	0
1	-0.031230	0.023416	0.008874	0.074721	-0.016033	-0.025616	-0.050069	-0.018542	0.036259	-0.002108	...	-0.076787	0.021046	-0.026546	0.046404	-0.064159	0.009849	0.009103	0.022082	0.063062	1

2 rows × 201 columns

[30] df_lemma.head(2)

	q1	q2	is_duplicate
48958	job microsoft certification	microsoft certification like mcase mc aid get L...	0
77587	write article art gallery	write art easy	0

lemma_df.head(2)

	0_left	1	2	3	4	5	6	7	8	9	...	41	42	43	44	45	46	47	48	49	0_right
0	-0.250281	0.268393	-0.507923	-0.090775	-0.015231	-0.105132	-0.230446	0.054211	0.073804	-0.101047	...	0.015447	-0.113064	0.021469	0.100241	-0.08906	0.092462	0.195167	0.278540	-0.025395	0
1	0.042417	0.109407	0.018437	-0.006771	-0.064517	-0.041608	-0.039554	-0.070556	0.149302	-0.058216	...	0.078449	0.072881	-0.073291	-0.031450	0.07713	0.034200	0.182466	0.055447	0.062686	0

2 rows × 101 columns

RANDOM FOREST

```
#splits = ShuffleSplit(n_splits = 1, test_size = .2, random_state = 0)

rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {'max_depth': [20, None],
                 'min_samples_split': [20]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resources=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, 'f1acc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

[75] #Random Forest

```
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)
```

RandomForestClassifier(max_depth=20, min_samples_split=20, n_estimators=5, random_state=0)

Classification Report - RF

	precision	recall	f1-score	support
0	1.00	1.00	1.00	4748
1	0.99	1.00	1.00	2752
accuracy	1.00	1.00	1.00	7500
macro avg	1.00	1.00	1.00	7500
weighted avg	1.00	1.00	1.00	7500

Confusion Matrix - RF

```
[[4727  21]
 [ 1 2751]]
```

Executing (54): Cell - fit() : _run_search() : evaluate_candidates() : _call_() : retrieve() : wrap_future_result() : result() : wait()

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = { 'max_depth': [5],
                 'min_samples_split': [20]
                 }

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resources=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, 'f1acc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

#Random Forest

```
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)
```

RandomForestClassifier(max_depth=5, min_samples_split=20, n_estimators=5, random_state=0)

Classification Report - RF

	precision	recall	f1-score	support
0	1.00	0.59	0.82	6883
1	0.25	1.00	0.40	697
accuracy	0.63	0.85	0.72	7580
macro avg	0.63	0.85	0.61	7580
weighted avg	0.93	0.72	0.78	7580

Confusion Matrix - RF

```
[[4728 2075]
 [ 0 697]]
```

6s - completed at 14:54

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = { "max_depth": [5],
                  "min_samples_split": [100]
                }

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                               max_resources=5,
                               random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_best, 'macro', rf_acc_best)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(max_depth=5, min_samples_split=100, n_estimators=5,
                       random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0         1.00      0.69      0.81      6890
1         0.22      1.00      0.36       610

accuracy          0.51      0.54      0.71      7500
macro avg          0.54      0.71      0.78      7500
weighted avg          0.54      0.71      0.78      7500

Confusion Matrix - RF
[[4728 2162]
 [   0   610]]
```

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = { "max_depth": [4],
                  "min_samples_split": [100]
                }

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                               max_resources=5,
                               random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_best, 'macro', rf_acc_best)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(max_depth=4, min_samples_split=100, n_estimators=5,
                       random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0         1.00      0.88      0.94      5381
1         0.76      1.00      0.87      2119

accuracy          0.88      0.54      0.91      7500
macro avg          0.88      0.54      0.90      7500
weighted avg          0.93      0.91      0.92      7500

Confusion Matrix - RF
[[4728  653]
 [   0 2119]]
```

SELECTED HP

```
def rf_clf(X_train, X_test, y_train, y_test):
    """Parameter tuning for Random Forest Classifier and model fit"""
    #split = ShuffleSplit(n_splits = 1, test_size = .1, random_state = 0)
    rf_clf = RandomForestClassifier(random_state = 0)
    rf_param_grid = { "max_depth": [1, 100],
                     "min_samples_split": [10]
                   }

    rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators', max_resources=5,
                                   random_state=0).fit(X_train, y_train)

    rf_model = rf_search.best_estimator_
    y_pred = rf_model.predict(X_test)

    log_loss = metrics.log_loss(y_test, y_pred)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    precision = metrics.precision_score(y_test, y_pred)
    recall = metrics.recall_score(y_test, y_pred)
    f1_score = metrics.f1_score(y_test, y_pred)
    print(rf_model)
    #print('log_loss', rf_loss_best, 'macro', rf_acc_best)
    return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model
```

SVC

```
splits = ShuffleSplit(n_splits = 1, test_size = .2, random_state = 0)

#svc_param_grid = {'C': [0.1, 10, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'linear', 'poly', 'rbf']}
#svc_param_grid = {'C': [0.1, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'poly', 'rbf']}
svc_param_grid = {'C': [10],
                  'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                  'kernel': ['sigmoid']}
svc_clf = SVC()
# gamma='scale'
svc_search = HalvingGridSearchCV(svc_clf, svc_param_grid, factor = 2, scoring = 'accuracy')
svc_search.fit(X_train, y_train)

svc_model = svc_search.best_params_
svc_model_est = svc_search.best_estimator_

y_pred = svc_model_est.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred, average='weighted')
recall = metrics.recall_score(y_test, y_pred, average='weighted')
f1_score = metrics.f1_score(y_test, y_pred, average='macro')
print(svc_model)
return accuracy, precision, recall, f1_score, log_loss, y_pred, svc_model_est

## SVC - using param grid - 100, 30, 1
acc_svc_100ma, precision_svc_100ma, recall_svc_100ma, f1_score_svc_100ma, log_loss_svc_100ma, y_pred_svc_100ma, svc_model_100ma = svc_clf(X_train_100ma, X_test_100ma, y_train_100ma, y_test_100ma)

report_svc_100ma = classification_report(y_pred_svc_100ma, y_test_100ma)
matrix_svc_100ma = confusion_matrix(y_pred_svc_100ma, y_test_100ma)

print('Classification Report - SVC\n', report_svc_100ma)
print()
print('Confusion Matrix - SVC\n', matrix_svc_100ma)

{'C': 1, 'gamma': 0.01, 'kernel': 'rbf'}
Classification Report - SVC
precision    recall  f1-score   support

0         1.00      1.00      1.00       4733
1         1.00      1.00      1.00       2767

accuracy          1.00          1.00          1.00       7500
macro avg          1.00          1.00          1.00       7500
weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - SVC
[[4733   0]
 [   0 2767]]
```

SELECTED

```
def svc_clf(X_train, X_test, y_train, y_test):
    """Parameter tuning for svm classifier and model fit"""

    splits = ShuffleSplit(n_splits = 1, test_size = .2, random_state = 0)

    #svc_param_grid = {'C': [0.1, 10, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'linear', 'poly', 'rbf']}
    #svc_param_grid = {'C': [0.1, 100], 'gamma': [1, 0.1, 0.01], 'kernel': ['sigmoid', 'poly', 'rbf']}
    svc_param_grid = {'C': [10],
                      'gamma': [1],
                      'kernel': ['sigmoid']}
    svc_clf = SVC()
    # gamma='scale'
    svc_search = HalvingGridSearchCV(svc_clf, svc_param_grid, factor = 2, scoring = 'accuracy')
    svc_search.fit(X_train, y_train)

    svc_model = svc_search.best_params_
    svc_model_est = svc_search.best_estimator_

    y_pred = svc_model_est.predict(X_test)

    log_loss = metrics.log_loss(y_test, y_pred)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    precision = metrics.precision_score(y_test, y_pred, average='weighted')
    recall = metrics.recall_score(y_test, y_pred, average='weighted')
    f1_score = metrics.f1_score(y_test, y_pred, average='macro')
    print(svc_model)
    return accuracy, precision, recall, f1_score, log_loss, y_pred, svc_model_est
```

DOC2VEC – RANDOM FOREST

VECTOR 20

```
#Random forest # doc0, window=1, vector=20, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       1.00      1.00      1.00       4737
    1       1.00      1.00      1.00       2763

 accuracy          1.00          1.00          1.00       7500
 macro avg          1.00          1.00          1.00       7500
 weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - RF
[[4732  5]
 [ 1 2762]]

[149] #Random forest # doc1, window=1, vector=20, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=20, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       1.00      1.00      1.00       4746
    1       0.99      1.00      1.00       2754

 accuracy          1.00          1.00          1.00       7500
 macro avg          1.00          1.00          1.00       7500
 weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - RF
[[4732 14]
 [ 1 2753]]
```

```
[ ] #Random forest # doc0, window=1, vector=100, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=20, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.99      0.98      0.98       4785
    1       0.96      0.98      0.97       2715

 accuracy          0.97          0.98          0.98       7500
 macro avg          0.97          0.98          0.98       7500
 weighted avg          0.98          0.98          0.98       7500

Confusion Matrix - RF
[[4674 111]
 [ 59 2656]]

[ ] #Random forest # doc0, window=1, vector=20, epochs=50
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       1.00      1.00      1.00       4737
    1       1.00      1.00      1.00       2763

 accuracy          1.00          1.00          1.00       7500
 macro avg          1.00          1.00          1.00       7500
 weighted avg          1.00          1.00          1.00       7500

Confusion Matrix - RF
[[4732  5]
 [ 1 2762]]
```

DM = 1 - VECTOR=50 – WINDOW= 3 - EPOCHS=30

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {"max_depth": [10, None],
                 "min_samples_split": [2, 4]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy',
                                max_resources=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print(log_loss, rf_loss_bow, "\nacc", rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=4, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.99      0.98      0.98      4763
1           0.96      0.98      0.97      2737

accuracy          0.98      0.98      7500
macro avg         0.97      0.98      0.98      7500
weighted avg      0.98      0.98      0.98      7500

Confusion Matrix - RF
[[4661 182]
 [ 62 2670]]
```

21s completed at 11:16

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {"max_depth": [10, None],
                 "min_samples_split": [20, 40]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy',
                                max_resources=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_
y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print(log_loss, rf_loss_bow, "\nacc", rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemna, precision_rf_lemna, recall_rf_lemna, f1_score_rf_lemna, log_loss_rf_lemna, y_pred_rf_lemna, berf_lemna = rf_clf(X_train_lemna, X_test_lemna, y_train_lemna, y_test_lemna)

report_rf_lemna = classification_report(y_pred_rf_lemna, y_test_lemna)
matrix_rf_lemna = confusion_matrix(y_pred_rf_lemna, y_test_lemna)
print('Classification Report - RF\n', report_rf_lemna)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemna)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0           0.98      0.95      0.97      4880
1           0.92      0.97      0.94      2620

accuracy          0.96      0.96      7500
macro avg         0.95      0.96      0.95      7500
weighted avg      0.96      0.96      0.96      7500

Confusion Matrix - RF
[[4646 234]
 [ 82 2538]]
```

16s completed at 11:17

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {'max_depth': [40, None],
                 "min_samples_split": [20, 40]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resource=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_

y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, '\nacc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0     0.98     0.95     0.97     4880
1     0.92     0.97     0.94     2620

accuracy          0.96     0.96     0.96     7500
macro avg         0.95     0.96     0.95     7500
weighted avg      0.96     0.96     0.96     7500

Confusion Matrix - RF
[[4646  234]
 [  82 2538]]
```

15s completed at 11:18

```
rf_clf = RandomForestClassifier(random_state = 0)
rf_param_grid = {'max_depth': [40, None],
                 "min_samples_split": [20, 40]}

rf_search = HalvingGridSearchCV(rf_clf, rf_param_grid, scoring='accuracy', resource='n_estimators',
                                max_resource=5,
                                random_state=0).fit(X_train, y_train)

rf_model = rf_search.best_estimator_

y_pred = rf_model.predict(X_test)

log_loss = metrics.log_loss(y_test, y_pred)
accuracy = metrics.accuracy_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)
recall = metrics.recall_score(y_test, y_pred)
f1_score = metrics.f1_score(y_test, y_pred)
print(rf_model)
#print('log_loss', rf_loss_bow, '\nacc', rf_acc_bow)
return accuracy, precision, recall, f1_score, log_loss, y_pred, rf_model

#Random Forest
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(max_depth=40, min_samples_split=20, n_estimators=3,
                        random_state=0)
Classification Report - RF
precision    recall  f1-score   support

0     0.99     0.98     0.99     4789
1     0.97     0.99     0.98     2711

accuracy          0.98     0.98     0.98     7500
macro avg         0.98     0.98     0.98     7500
weighted avg      0.98     0.98     0.98     7500

Confusion Matrix - RF
[[4696   93]
 [  32 2679]]
```

16s completed at 11:19

VECTOR 100 VS 200 (DM = 1 – WINDOW = 3, EPOCHS = 30)

```
#Random Forest # dm1, window3, vector=200, epochs=30
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.96       0.91       0.94       5002
    1       0.84       0.93       0.89       2498

 accuracy      0.90      0.92      0.91       7500
 macro avg      0.90      0.92      0.91       7500
 weighted avg      0.92      0.92      0.92       7500

Confusion Matrix - RF
[[4565 437]
 [ 108 2330]]
```

CHANGING NUMBER OF WINDOW

```
[37] #Random Forest # 100 - 30 - 1
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.99       0.96       0.98       4865
    1       0.93       0.98       0.96       2635

 accuracy      0.96      0.97      0.97       7500
 macro avg      0.96      0.97      0.97       7500
 weighted avg      0.97      0.97      0.97       7500

Confusion Matrix - RF
[[4680 185]
 [   53 2520]]
```

```
#Random Forest # 100 - 30 - 3
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=20, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.96       0.93       0.95       4867
    1       0.88       0.93       0.91       2633

 accuracy      0.92      0.93      0.93       7500
 macro avg      0.92      0.93      0.93       7500
 weighted avg      0.93      0.93      0.93       7500

Confusion Matrix - RF
[[4548 319]
 [ 191 2442]]
```

```
#Random Forest # 100 - 30 - 5
acc_rf_lemma, precision_rf_lemma, recall_rf_lemma, f1_score_rf_lemma, log_loss_rf_lemma, y_pred_rf_lemma, berrf_lemma = rf_clf(X_train_lemma, X_test_lemma, y_train_lemma, y_test_lemma)

report_rf_lemma = classification_report(y_pred_rf_lemma, y_test_lemma)
matrix_rf_lemma = confusion_matrix(y_pred_rf_lemma, y_test_lemma)
print('Classification Report - RF\n', report_rf_lemma)
print()
print('Confusion Matrix - RF\n', matrix_rf_lemma)

RandomForestClassifier(min_samples_split=40, n_estimators=3, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

    0       0.99       0.95       0.97       4915
    1       0.91       0.98       0.95       2585

 accuracy      0.95      0.97      0.96       7500
 macro avg      0.95      0.97      0.96       7500
 weighted avg      0.96      0.96      0.96       7500

Confusion Matrix - RF
[[4679 236]
 [   54 2521]]
```

RESULTS

Lemmatized without stopwords

Not tokenized – 100 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
              precision    recall  f1-score   support

    0       1.00      0.90      0.95        150
    1       0.93      1.00      0.96        104

 accuracy      0.95      0.95      0.97        300
 macro avg     0.95      0.95      0.97        300
 weighted avg   0.95      0.97      0.97        300

Confusion Matrix - RF
[[150  0]
 [ 0 104]]
{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
              precision    recall  f1-score   support

    0       0.76      0.77      0.77        185
    1       0.62      0.61      0.62        115

 accuracy      0.69      0.71      0.71        300
 macro avg     0.69      0.69      0.69        300
 weighted avg   0.71      0.71      0.71        300

Confusion Matrix - SVC
[[145  40]
 [ 40 155]]
```

Tokenized – 100 VECTORS

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
              precision    recall  f1-score   support

    0       0.97      0.94      0.96        485
    1       0.90      0.95      0.92        265

 accuracy      0.94      0.95      0.95        750
 macro avg     0.94      0.95      0.94        750
 weighted avg   0.95      0.95      0.95        750

Confusion Matrix - RF
[[487  63]
 [ 63 263]]
{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
              precision    recall  f1-score   support

    0       0.78      0.76      0.77        485
    1       0.58      0.61      0.59        265

 accuracy      0.68      0.68      0.68        750
 macro avg     0.68      0.68      0.68        750
 weighted avg   0.71      0.71      0.71        750

Confusion Matrix - SVC
[[367 118]
 [118 383]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
              precision    recall  f1-score   support

    0       0.88      0.88      0.88        475
    1       0.79      0.81      0.80        275

 accuracy      0.84      0.85      0.85        750
 macro avg     0.84      0.85      0.85        750
 weighted avg   0.85      0.85      0.85        750

Confusion Matrix - RF
[[475  75]
 [ 75 275]]
{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
              precision    recall  f1-score   support

    0       0.79      0.78      0.78        475
    1       0.62      0.64      0.63        275

 accuracy      0.78      0.71      0.75        750
 macro avg     0.78      0.71      0.75        750
 weighted avg   0.73      0.73      0.73        750

Confusion Matrix - SVC
[[371 104]
 [104 379]]
```


Lemmatized with stopwords

Not tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall    f1-score   support

 0 1.00 0.57 0.70 103
 1 0.96 1.00 0.98 107

 accuracy 0.98
 macro avg 0.98 0.59 0.98
 weighted avg 0.98 0.98 0.98

Confusion Matrix - RF
[[100  3]
 [ 0 107]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall    f1-score   support

 0 0.74 0.74 0.74 103
 1 0.56 0.57 0.57 111

 accuracy 0.65
 macro avg 0.65 0.65 0.65
 weighted avg 0.62 0.62 0.62

Confusion Matrix - SVC
[[100  3]
 [ 0 111]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall    f1-score   support

 0 0.97 0.92 0.95 405
 1 0.86 0.95 0.90 255

 accuracy 0.92
 macro avg 0.92 0.93 0.92
 weighted avg 0.93 0.93 0.93

Confusion Matrix - RF
[[406 39]
 [ 0 255]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall    f1-score   support

 0 0.72 0.71 0.72 472
 1 0.52 0.52 0.52 178

 accuracy 0.62
 macro avg 0.62 0.62 0.62
 weighted avg 0.64 0.64 0.64

Confusion Matrix - SVC
[[137 15]
 [13 163]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall    f1-score   support

 0 0.94 0.86 0.90 513
 1 0.79 0.88 0.81 237

 accuracy 0.87
 macro avg 0.84 0.87 0.85
 weighted avg 0.88 0.87 0.87

Confusion Matrix - RF
[[442 71]
 [ 0 237]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall    f1-score   support

 0 0.75 0.74 0.74 475
 1 0.56 0.57 0.56 275

 accuracy 0.68
 macro avg 0.68 0.68 0.68
 weighted avg 0.68 0.68 0.68

Confusion Matrix - SVC
[[183 124]
 [19 194]]
```

With stopwords, not lemmatized nor stemmed

Not tokenized – 100 vectors

```
RandomForestClassifier(m_n_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.99      0.96      0.98      193
     1       0.94      0.90      0.92      257

 accuracy      0.96
 macro avg     0.96      0.93      0.95      450
 weighted avg   0.97      0.93      0.95      450

Confusion Matrix - RF
[[186  2]
 [ 3 267]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.77      0.77      0.77      187
     1       0.62      0.61      0.61      113

 accuracy      0.69
 macro avg     0.69      0.69      0.69      300
 weighted avg   0.71      0.71      0.71      300

Confusion Matrix - SVC
[[144 43]
 [ 44 193]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(m_n_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.97      0.96      0.96      483
     1       0.96      0.98      0.97      267

 accuracy      0.96
 macro avg     0.96      0.97      0.97      750
 weighted avg   0.96      0.96      0.96      750

Confusion Matrix - RF
[[480  3]
 [ 15 267]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.69      0.69      0.69      478
     1       0.49      0.49      0.49      268

 accuracy      0.63
 macro avg     0.63      0.63      0.63      746
 weighted avg   0.62      0.62      0.62      746

Confusion Matrix - SVC
[[126 144]
 [144 194]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(m_n_samples_split=4, n_estimators=5, random_state=0)
Classification report - RF
      precision    recall  f1-score   support

     0       0.97      0.93      0.95      467
     1       0.88      0.94      0.91      283

 accuracy      0.93
 macro avg     0.93      0.94      0.93      750
 weighted avg   0.94      0.93      0.94      750

Confusion Matrix - RF
[[454 33]
 [ 36 287]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
      precision    recall  f1-score   support

     0       0.70      0.69      0.70      473
     1       0.48      0.50      0.49      277

 accuracy      0.69
 macro avg     0.69      0.69      0.69      750
 weighted avg   0.62      0.62      0.62      750

Confusion Matrix - SVC
[[129 145]
 [142 193]]
```

Without stopwords, not lemmatized nor stemmed

No tokenized – 100 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=4)
Classification Report - RF
      precision    recall  f1-score   support

     0       1.00      0.84      0.91       223
     1       0.69      1.00      0.83        77

 accuracy      0.84      0.92      0.86       300
 macro avg      0.84      0.92      0.86       300
 weighted avg      0.92      0.88      0.89       300

Confusion Matrix - RF
[[188  35]
 [  9  77]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.82      0.78      0.80       188
     1       0.62      0.68      0.64       112

 accuracy      0.73      0.73      0.73       300
 macro avg      0.72      0.73      0.72       300
 weighted avg      0.75      0.75      0.75       300

Confusion Matrix - SVC
[[155  43]
 [ 33  89]]
```

Tokenized - vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=4)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.97      0.97      0.97       471
     1       0.95      0.96      0.96       279

 accuracy      0.96      0.96      0.96       750
 macro avg      0.96      0.96      0.96       750
 weighted avg      0.97      0.97      0.97       750

Confusion Matrix - RF
[[466  25]
 [ 25  267]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.79      0.78      0.78       474
     1       0.62      0.63      0.63       276

 accuracy      0.71      0.72      0.72       750
 macro avg      0.71      0.71      0.71       750
 weighted avg      0.73      0.73      0.73       750

Confusion Matrix - SVC
[[369 186]
 [186 171]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(min_samples_split=4, n_estimators=5, random_state=4)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.96      0.98      0.97       468
     1       0.82      0.92      0.87       282

 accuracy      0.89      0.95      0.92       750
 macro avg      0.89      0.95      0.92       750
 weighted avg      0.91      0.95      0.93       750

Confusion Matrix - RF
[[461  47]
 [ 29 281]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.81      0.79      0.80       478
     1       0.65      0.67      0.66       272

 accuracy      0.73      0.73      0.73       750
 macro avg      0.73      0.73      0.73       750
 weighted avg      0.75      0.75      0.75       750

Confusion Matrix - SVC
[[379  89]
 [ 81 181]]
```

Stemmed with stopwords

Not tokenized – 100 vectors

```
RandomForestClassifier(m_n_samples_split=40, n_estimators=5, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.96      0.91      0.93       180
     1       0.84      0.70      0.76       182

 accuracy      0.90      0.82      0.91       360
 macro avg      0.90      0.82      0.91       360
 weighted avg      0.92      0.71      0.91       360

Confusion Matrix - RF
[[180  31]
 [ 30 181]]

{'C': 18, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.91      0.85      0.88       203
     1       0.72      0.84      0.78       97

 accuracy      0.86      0.84      0.85       300
 macro avg      0.82      0.84      0.83       300
 weighted avg      0.85      0.84      0.85       300

Confusion Matrix - SVC
[[172  31]
 [ 30 117]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(m_n_samples_split=40, n_estimators=5, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.98      0.95      0.97       488
     1       0.91      0.97      0.94       264

 accuracy      0.95      0.96      0.96       752
 macro avg      0.95      0.96      0.96       752
 weighted avg      0.96      0.96      0.96       752

Confusion Matrix - RF
[[482  34]
 [ 30 264]]

{'C': 18, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.75      0.72      0.74       482
     1       0.53      0.56      0.54       264

 accuracy      0.64      0.64      0.64       752
 macro avg      0.64      0.64      0.64       752
 weighted avg      0.67      0.67      0.67       752

Confusion Matrix - SVC
[[391 181]
 [119 189]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(m_n_samples_split=40, n_estimators=5, random_state=0)
Classification Report - RF
      precision    recall  f1-score   support

     0       0.94      0.87      0.91       518
     1       0.76      0.89      0.82       240

 accuracy      0.85      0.88      0.86       758
 macro avg      0.85      0.88      0.86       758
 weighted avg      0.89      0.88      0.88       758

Confusion Matrix - RF
[[444  64]
 [ 30 244]]

{'C': 18, 'gamma': 0, 'kernel': 'sigmoid'}
Classification Report - SVC
      precision    recall  f1-score   support

     0       0.76      0.74      0.75       482
     1       0.56      0.59      0.57       268

 accuracy      0.69      0.69      0.69       758
 macro avg      0.66      0.67      0.66       758
 weighted avg      0.69      0.69      0.69       758

Confusion Matrix - SVC
[[399 213]
 [111 157]]
```

Stemmed without stopwords

No Tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

0       0.97      0.93      0.95      196
1       0.88      0.95      0.92      184

accuracy          0.93
macro avg          0.93
weighted avg       0.94

Confusion Matrix - RF
[[183  11]
 [ 9 181]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

0       0.95      0.85      0.90      200
1       0.72      0.89      0.80      91

accuracy          0.81
macro avg          0.83
weighted avg       0.85

Confusion Matrix - SVC
[[176  24]
 [ 18  81]]
```

Tokenized – 100 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

0       0.98      0.97      0.98      471
1       0.95      0.96      0.96      277

accuracy          0.97
macro avg          0.97
weighted avg       0.97

Confusion Matrix - RF
[[466  5]
 [ 18 267]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

0       0.72      0.72      0.72      471
1       0.53      0.53      0.53      279

accuracy          0.65
macro avg          0.62
weighted avg       0.65

Confusion Matrix - SVC
[[339 132]
 [119 140]]
```

Tokenized – 200 vectors

```
RandomForestClassifier(n_estimators=5, random_state=0)
Classification report - RF
precision    recall  f1-score   support

0       0.93      0.88      0.90      497
1       0.78      0.86      0.82      253

accuracy          0.87
macro avg          0.85
weighted avg       0.87

Confusion Matrix - RF
[[459  38]
 [ 16 218]]

{'C': 10, 'gamma': 0, 'kernel': 'sigmoid'}
Classification report - SVC
precision    recall  f1-score   support

0       0.71      0.72      0.71      462
1       0.54      0.52      0.53      288

accuracy          0.62
macro avg          0.62
weighted avg       0.64

Confusion Matrix - SVC
[[332 130]
 [128 136]]
```